

Attention Research

Results across two projects

PROJECT 1

CRPA: is behavioural contribution measurable at all?

Audited a partitioned-attention repository, then tested whether its own central claim survives a resolvability check.

1.1 The claim that was withdrawn

The branch reported that structural overlap does not predict behavioural contribution. Check 0 measures each edge's effect twice on disjoint halves of the evaluation split and correlates the two estimates. **Pooled verdict: UNRESOLVABLE.**

seed	r_delta	r_stat	ceiling	observed rho	delta size	converges
42	+0.088	+0.119	0.102	+0.018	6.0 ULP	no
1337	-0.026	+0.255	0.000	-0.017	5.0 ULP	no
2024	+0.012	+0.250	0.054	+0.013	4.0 ULP	no




Figure 1. Observed correlation against the ceiling that measurement unreliability imposes on it, $\sqrt{r_{\text{delta}} \times r_{\text{stat}}}$. Every observed value sits far inside its own ceiling, so a real decoupling and measured noise cannot be told apart. Seed 1337 has no bar: its split-half reliability is negative, so the ceiling degenerates to zero and that seed carries no information at all. $n = 3$ seeds, 1,044 to 1,152 scored edges each.

1.2 The floor holds across three orders of magnitude

The replacement claim is that single-edge contribution is not resolvable, not that overlap fails to predict it.

scale	precision	single-edge delta
12.4M (Tier 1)	float32	4 to 6 ULP, split-half reliability 0.088
138M (Tier 2)	float32	exactly one ULP, 9.5367e-07
6.9B (Tier 3)	bfloat16	exactly 0.000e+00 across 192 edges




Figure 2. Typical single-edge intervention effect, expressed in units of the least representable step at the working precision. At 6.9B the effect is not small but exactly zero: the edit shifts logits by 2.58, 3.05 and 0.127 at layers 0, 16 and 31, and the loss is bit-identical to baseline for all 192 edges. A criterion thresholding this quantity is unfalsifiable at every scale tested.

1.3 Reproduction of the published table

Re-run from the original commit. Four of five rows come back. The one that does not is the one the paper's claim rests on.

Variant	Published	Reproduced	Verdict
Dense Transformer	50.9%	55.6%	reproduces
Sliding Window	51.9%	52.5%	reproduces
CRPA no regularisation	8.4%	10.9%	reproduces
CRPA naive regularisation	5.3%	5.3%	reproduces exactly
CRPA causal regularisation	32.8%	7.5%	does not reproduce

1.4 Three seeds against a measured chance floor

The floor is measured by simulating the task generator, not asserted from vocabulary size: 52.78% (sd 0.08, n = 25,000 x 5). **No variant clears it, dense included.** The retrieval task cannot discriminate between these variants.

Variant	Retrieval	t vs floor	Verdict	Overlap
dense	53.6%	+0.8	indistinguishable	0.403
sliding window	47.2%	-1.4	indistinguishable	0.322
crpa_noreg	4.4%	-66	below	0.244
crpa_naive	4.6%	-100	below	0.264
crpa_contribution	4.3%	-50	below	0.222

1.5 Long context: measured cost, and a proven wall

A fused dense kernel is faster at every length measured, at equal memory, with the gap narrowing as context grows.

Context	Dense	Sliding	CRPA	CRPA vs dense	Peak memory
4,096	13.2 ms	19.5 ms	31.5 ms	2.39x slower	737 MB
8,192	30.0 ms	62.4 ms	60.1 ms	2.00x slower	1,181 MB
16,384	72.0 ms	218.1 ms	119.1 ms	1.65x slower	2,067 MB
32,768	–	–	–	<i>not run</i>	<i>out of memory</i>
65,536	–	–	–	<i>not run</i>	<i>out of memory</i>

32k and 64k were retried five times, including on an idle A100 80GB at the correct profile with batch size 1 and adaptive chunking. All exceed memory. The investigation established something more useful than the failure: **a forward pass at 32k costs 1.90 GB peak**, so the model scales and the cost sits in the diagnostic, not in inference.




Figure 3. Measured forward latency and peak allocated memory at the three context lengths that fit. A fused dense kernel is faster than the sparse gather path at every length, and the gap narrows from 2.39x to 1.65x as context grows, so the sparse path is closing but has not overtaken it by 16k. The shaded region marks lengths that exhausted memory on every attempt and are reported as not run rather than extrapolated. Batch size 1, bfloat16, 138M parameters, median of 10 iterations after 5 warmups.

1.6 Defects found and repaired

Each of these changed a reported number.

Defect	Consequence
Central claim unsupported by its own measurement	A decoupling claim the data cannot support. Withdrawn
Chance floor asserted, not measured	Real floor 52.78%, not 5.0%. Every above-chance verdict inverted
Intervention masked a query-row pair, not a query-to-key edge	The measured object was not the object being reasoned about
Roughly half of all interventions were causally inert	Delta zero by construction, so the pair was filed redundant
Short runs recordable as completed measurements	Nine 3-iteration runs entered an aggregate and moved it by 26x
Auxiliary routing entropy pinned at $\ln(4)$	A constant reported as a finding

xsa-controls: sanity checks for attention surgery

Three checks shipped as importable code: is the measurement resolvable, does the statistic beat an anisotropy null, and does a matched arbitrary intervention recover the gain?

2.1 Check 1 across a nine-model scale ladder

5,408 head-level rows, 32 real documents per model, eager attention, null partner drawn within the sequence from positions the query could causally attend.

model	params	cos_self	cos_null	excess	% self-specific
gpt2	124M	0.4828	0.2987	0.1840	38.1
pythia-160m	160M	0.4180	0.2637	0.1544	36.9
gpt2-medium	355M	0.4252	0.2579	0.1674	39.4
pythia-410m	410M	0.4022	0.1937	0.2086	51.9
gpt2-large	774M	0.4213	0.2117	0.2096	49.7
pythia-1.4b	1.4B	0.3862	0.1900	0.1963	50.8
gpt2-xl	1.5B	0.3861	0.2069	0.1792	46.4
pythia-2.8b	2.8B	0.3565	0.1859	0.1705	47.8
pythia-6.9b	6.9B	0.3404	0.1979	0.1425	41.9




Figure 4. The anisotropy null weakens with scale but does not vanish. cos_null falls from 0.2637 at 160M to 0.1979 at 6.9B, yet 58% of the raw statistic is still explained by the null at 6.9B, and about half across the shaded range where the method was actually trained. This answers the strongest objection to the framing, that large models are isotropic, with measurement rather than extrapolation. $n = 9$ models, 5,408 heads.

2.2 The null depends on where it is measured

The same model and corpus, measured at five context lengths.

Context length	64	128	256	512	1024
$\cos(y, v)$ observed	0.4900	0.4846	0.4812	0.4772	0.4743
$\cos(y, v')$ null	0.3747	0.3361	0.3096	0.2904	0.2788
self-specific share	23.5%	30.7%	35.7%	39.2%	41.2%




Figure 5. The observed statistic is nearly flat across a sixteenfold change in context length while the null moves by a third, so the self-specific share swings from 23.5% to 41.2%. A claim of the form “only N% of this statistic is self-specific” is therefore a property of the measurement context as much as of the model, and Check 1 must always be reported with the length it was measured at. $n = 24$ documents per length, GPT-2.

2.3 Grouped-query attention, which nobody had checked

Under GQA several query heads share one KV head, so the token's own value vector is shared across a group. The source method's own table reports no KV-head count.

model	query / KV heads	within-group excess	across-group excess
Qwen2.5-0.5B	14 / 2	+0.2415	-0.1876
Qwen2.5-1.5B	12 / 2	+0.2731	-0.1922




Figure 6. Borrowing a neighbouring KV group's value at the same position does not merely lose the effect, it reverses it. The self-value direction is specific to the head's own group. GQA models also show a higher self-specific share, 56 to 59%, than any multi-head model on the ladder, 37 to 52%, so grouped-query attention is not a neutral change of variable for this family of methods. $n = 2$ models, 672 heads.

2.4 The matched intervention, at full seed count

24 cells, 3 arms x 8 seeds, identical initialisation and data order per seed. Measured step-0 deviation across all five arms is 0.000e+00, so the arms start from a common point.

arm	mean delta vs baseline	95% CI	t	p	n
random (primary)	+0.001190	[+0.000351, +0.002040]	+2.48	0.042	8
xsa	+0.001515	[-0.001223, +0.004807]	+0.92	0.387	8




Figure 7. Both arms fall inside the shaded band of effects the design cannot resolve. The minimum detectable effect is 0.00518 nats against the 0.00076 the method's own independent replication reports, so the study is underpowered about sevenfold and cannot settle the question in either direction. The two arms are also indistinguishable from each other, which is what Check 2 asks. Reported as a power failure, not as a null result. $n = 8$ seeds per arm.

2.5 The checklist discriminates between methods

Check 1 applied to two further methods on frozen GPT-2, each with a null matched to the structure the statistic inherits for free.

method	statistic	null	excess	% self-specific	survives
attention sinks	0.4028	0.0034	0.3994	99.2	yes
massive activations	12.8719	3.6452	9.2266	71.7	yes
self-value (this work, 6.9B)	0.3404	0.1979	0.1425	41.9	no




Figure 8. Attention sinks retain 99.2% of their statistic after a null matched for recency, and massive activations 71.7% against the maximum expected from a matched-variance Gaussian. The self-value statistic retains 41.9%. The checklist separates methods rather than debunking all of them, which is a stronger position than a blanket null.

2.6 A reproduction left failing on purpose

Published GPT-2 reference: `cos_self` 0.5406, `cos_null` 0.3798, excess 0.1608. Thirteen measurement conventions were tested and the full grid is reported unselected.

convention varied	range of <code>cos_self</code> observed	reproduces reference
sequence length, 64 to 1024	0.4827 to 0.4955	no
position 0 included	0.4847	no
null partner definition, 3 variants	0.4832	no
head pooling	0.4832	no
layer subsets, 4 variants	0.4405 to 0.6494	no

None of the thirteen reproduces the reference. Layer subsets move the statistic most, so a subset is the likeliest remaining explanation, but none lands on the reference triple. Reported as unexplained rather than resolved by further search: selecting the setting that hits a target and presenting it as the method is the practice this work argues against.

2.7 Defects found only by running it on GPU

None of these were visible from CPU testing.

Defect	Consequence
diagmask crashed on every GPU run	bf16 autocast dtype mismatch. One of five arms could never have run
Calibration output was never consumed	The factorial used a default budget, so the gate changed nothing
Budget clamped up to an unaffordable floor	Returned a plan that could not be paid for and called it fine
The spend threshold was never enforced	Defined in config, referenced nowhere
Two runs mixed token budgets in one directory	Caught by a guard ported from the sibling project; data discarded

Status, and what remains open

Both deliverables carry green CI. No compute is running.

	CRPA	xsa-controls
Tests	266 passing	202 passing
Coverage	83%	90%
Figures from real data	7 of 7	5 of 5
Compute used	A6000, RTX 6000 Ada, A100 80GB	RTX 6000 Ada, A100 80GB
Open item	long context at 32k and 64k	GPT-2 reference reproduction
Why it is open	proven memory wall over five attempts	thirteen conventions tested, none matches